EMOTION DETECTION AND SENTIMENT ANALYSIS

A PROJECT REPORT

Submitted by

Shourya Saxena(22BCE11051)
Harshvardhan Singh(22BCE10625)
Nipun Sharma(22BCE11190)
Anshika Mishra(22BCE10635)
Priyanshu Kumar(22BCE10623)

in partial fulfillment for the award of the degree of

BACHELOR OF TECHNOLOGY

in

COMPUTER SCIENCE AND ENGINEERING

SCHOOL OF COMPUTING SCIENCE AND ENGINEERING

VIT BHOPAL UNIVERSITY

KOTHRIKALAN, SEHORE MADHYA PRADESH – 466114

VIT BHOPAL UNIVERSITY, KOTHRIKALAN,

SEHORE

MADHYA PRADESH – 466114

BONAFIDE CERTIFICATE

Certified that this project report titled **“Emotion Detection and Sentiment Analysis”** is the bonafide work of

“Nipun Sharma(22BCE11190) Harshvardhan Singh(22BCE10625) Anshika Mishra (22BCE10635) Priyanshu Kumar(22BCE10623)”who carried out the project work under my supervision. Certified further that to the best of my knowledge the work reported at this time does not form part of any other project/research work based on which a degree or award was conferred on an earlier occasion on this or any other candidate.

PROGRAM CHAIR
Dr Preetam Suman
Program Chair, B.Tech CSE SCSE
VIT BHOPAL UNIVERSITY UNIVERSITY

PROJECT GUIDE
Dr. Vipin Jain
Professor,
VIT BHOPAL

The Project Exhibition II Examination is held on ______________

ACKNOWLEDGEMENT

First and foremost, I would like to thank the Lord Almighty for His presence and immense blessings throughout the project work.

I wish to express my heartfelt gratitude to Dr. S.Poonkuntran Professor and Dean, School of Computer Science and Engineering for much of his valuable support encouragement in carrying out this work.

I would like to thank my internal guide Dr. Vipin Jain , for continually guiding and actively participating in my project, giving valuable suggestions to complete the project work.

I would like to thank all the technical and teaching staff of the School of Computer Science and Engineering, who extended directly or indirectly all support.

Last, but not least, I am deeply indebted to my parents who have been the greatest support while I worked day and night for the project to make it a success

ABSTRACT

Understanding human emotions is a key part of social interaction and communication that is gaining more and more attention in the digital age. The proliferation of social media platforms has resulted in large amounts of data containing emotional signals that require the development of robust methods for emotion detection and analysis. This paper provides a comprehensive overview of current methods and developments in the field, with a particular focus on multimodal data integration and machine learning algorithms.

Emotion detection from facial expressions is of primary interest for this study. Facial expressions are considered a universal language that expresses the spectrum of human emotions. With the development of advanced algorithms, it has become easier to detect emotions from images automatically. The article discusses various methods and approaches used in facial emotion recognition, highlighting successes and challenges in achieving accurate and reliable results.

The development of the history of sentiment analysis provides valuable insight into the roots of the field. Previous efforts in sentiment analysis mainly focus on text data using vocabulary and machine learning techniques to classify text polarity. However, the emergence of multimodal data sources such as images, videos, and emoticons has prompted researchers to explore a more holistic approach to sentiment analysis. The combination of visual and textual information has led to significant progress in understanding the nuances of emotions from different perspectives.

Recent research in multimodal sentiment analysis focuses on the use of deep learning architecture and training methods. This approach allows combining visual and textual features that improve the interpretation and accuracy of sentiment analysis models. In addition, the article discusses the problem of linguistic errors, profanity, and ambiguity which is still an obstacle for emotional analysis.

In addition, the article presents applications for sentiment detection and sentiment analysis in various fields such as customer reanalysis. . , monitor political sentiment and monitor social media. By understanding user feelings and emotions, companies and organizations can make informed decisions to improve user experience and adjust their strategies.

Finally, the article highlights the importance of continuous research and innovation in emotion detection and emotion analysis. With the growth of social media, there is a demand for advanced algorithms that can understand human emotions in different contexts. The review provides a roadmap for future research directions and emphasizes the need for interdisciplinary collaboration and the integration of new technologies to advance emotion recognition and emotion analysis in the digital age.

Developing this topic allows us to go deeper into the complexity of understanding human emotions. . through digital data analysis. Emotions play an important role in shaping interpersonal communication, influencing the decision-making process, and managerial behavior. In today's world dominated by social media platforms, people express their feelings in various ways such as text, pictures, videos, animations and even respond to posts or comments. This diverse array of emotional responses presents opportunities and challenges for researchers seeking to develop effective methods for detecting and analyzing emotions.

Faces have long been considered powerful indicators of emotion in many cultures. The ability to automatically recognize emotions based on images opens up new application possibilities, for example in psychology, marketing, human-computer interaction, etc. by analyzing subtle changes in facial features such as eyebrow movement or curvature. lips, the algorithm can more accurately determine the emotional state of the situation.

Advances in sentiment analysis represent a move toward a more comprehensive approach that considers multiple types of data. . While previous methods focused on text analysis to determine the polarity of sentiments using dictionaries or machine learning models, modern methods incorporate visual information from images or videos to better understand sentiments expressed online. Emoticons have become valuable indicators of emotional tone in text-based communication, adding a new layer of complexity to the problem of emotional analysis.

Recent advances in deep learning have revolutionized multimodal sentiment analysis, enabling the exploration of complex patterns both visually and textually. simultaneous entry. Transfer learning techniques further improve model performance by using pre-trained representations from large databases. Researchers can effectively integrate information from different categories to capture feelings.

TABLE OF CONTENT

CHAPTER NO.	TITLE	PAGE NO.
	Abstract List Of Figures and Graphs List of Abbrevations	
1	**CHAPTER-1:** **PROJECT DESCRIPTION AND OUTLINE** 1.1 Introduction 1.2 Motivation for the work 1.3 Introduction to the model 1.4 Problem Statement 1.5 Objective	
2	**CHAPTER-2:** **LITERATURE SURVEY** 2.1 Outline 2.2 Core area of the project 2.3 Existing methods for Movie Recommendation	

<table>
<tr><td>3</td><td>CHAPTER-3:

PROJECT PROCEDURE

1. Data collection and preprocessing
2. Model selection and training
3. Sentiment Analysis
4. Evaluation and Analysis
5. Integration and Refinement
6. Testing and Deployment
7. Reporting and Documentation</td><td></td></tr>
<tr><td>4</td><td>CHAPTER-4:

WORK DONE

1. Training data
2. Creating Model
3. Configure Model
4. Train Model
5. Get predictions
6. Result</td><td></td></tr>
</table>

	CHAPTER-5: **OBSERVATION**	
	CHAPTER-6: **RESULT AND CONCLUSION**	
	CHAPTER-7: **RECOMMENDATIOIN FOR FUTURE WORK**	

	CHAPTER-8: **REFERENCES**	

LIST OF FIGURES AND GRAPHS

FIGURE NO.	FIGURE AND GRAPH	PAGE NO.
1	CNN – Flow Graph	
3(a)	JAFFE	
3(b)	Text Processing	
3(c)	RNN	
3(d)	LSTM	
3(e)	Bi-LSTM	
3(f)	Training Process	
4(a)	Custom CNN	

S.NO.	SYMBOL	**ABBREVATION**
1	CNN	Convolutional Neural Network
2	RNN	Recurrent Neural Network
3	RNDM	Recursive Neural Deep Model
4	NLP	Natural Language Processing
5	SVM	Support Vector Machine.
6	LDL	Label distribution learning
7	CPNN	Conditional Probability neural networks
8	BCPNN	Binary Conditional Probability neural networks
9	ACPNN	Augmented Conditional Probability neural networks
10	IAPS	International Affective Picture System
11	SVR	Support Vector Regression

12	FER	Facial emotion recognition
13	LSTM	Long Short-Term Memory
14	CBOWDA-LR	Continuous Bag of Words with Dimension Adaptive Learning Rate
15	Bi-LSTM	Bidirectional Long Short-Term Memory
16	T4SA	Twitter for Sentiment Analysis
17	FPGA	Field- Programmable Gate Array
18	GPU	Graphics Processing Units
19	DCNN	Dilated Convolutional Neural Network
20	MNB	Multinomial Naive Baye
21	NLTK	Natural Language Toolkit
22	POS	Part-of-speech
23	NAVA	Noun Adverb, verb and Adjective

CHAPTER 1

PROJECT DESCRIPTION AND OUTLINE

1. INTRODUCTION

The ability to understand human emotions and sentiments has long been a cornerstone of human interaction. This quest for deciphering emotions extends far back, with early philosophers like Aristotle grappling with the nature of human feeling. In the digital age, the explosion of social media platforms has created a vast trove of user-generated data, brimming with emotional cues. Sentiment analysis, the computational process of identifying the emotional tone of text, has emerged as a powerful tool for extracting insights from this data.

1.1 Tracing the Roots: A Historical Perspective on Sentiment Analysis

The origins of sentiment analysis can be traced back to the 1950s, where early efforts focused on analyzing written documents (Kdnuggets, 2015)[1]. These initial forays were primarily concerned with gauging overall sentiment, often employing rudimentary methods like identifying positive and negative words. The field received a significant boost with the advent of machine learning in the late 20th century. Algorithms began to analyze sentiments based on context, word placement, and the subtleties of language, leading to a more nuanced understanding of emotional expression (Archive App, 2021).

1.2 Building the Foundation: Early Work in Sentiment Analysis

The limitations of analyzing solely textual data became apparent as social media evolved into a more multimodal landscape. Images, videos, and emojis all convey emotional information that text alone might miss. Recognizing this potential, recent research has delved into multimodal sentiment analysis, which incorporates visual and auditory information alongside text (Clavel & Callejas, 2018)[2]. This approach offers a more comprehensive understanding of sentiment by leveraging the richness of social media data. For instance, facial expressions in videos or the color palette of an image can provide valuable clues about the underlying emotional tone.

1.3 The Evolving Landscape: Challenges and Opportunities

Despite the advancements, sentiment analysis still faces challenges. Understanding negation, sarcasm, and ambiguity in language remains an ongoing hurdle (MECS Press, 2013). Additionally, the ever-evolving nature of social media language necessitates continuous adaptation of sentiment analysis models. However, the potential benefits are vast. Businesses can gain invaluable insights into customer sentiment, political campaigns can gauge public opinion with greater accuracy, and social media platforms can identify and address potential issues like cyberbullying. This paper delves deeper into the world of emotion and sentiment analysis. Section 2 provides a comprehensive definition of sentiment analysis and emotion detection. Section 3 reviews existing sentiment analysis techniques, encompassing lexicon-based and machine learning approaches. Section 4 explores the emerging field of multimodal sentiment analysis, discussing its potential and applications. Section 5 delves into the challenges and opportunities associated with emotion and sentiment analysis in the age of social media. Finally, Section 6 concludes the paper by summarizing the key findings and outlining promising directions for future research.

1.2 MOTIVATION FOR THE WORK

Our project is driven by a fascination with the convergence of technology and human emotion. Inspired by the quest to decode sentiments, we aim to revolutionize user experiences in entertainment. Through advanced algorithms and data analytics, we seek to provide personalized movie recommendations, enriching engagement and fostering deeper connections between audiences and content. Our motivation lies in pioneering innovative approaches that push the boundaries of what's possible, with the goal of creating a more immersive and fulfilling entertainment landscape.

1.3 INTRODUCTION TO THE MODEL

This section introduces the concept of Convolutional Neural Networks (CNNs) in the context of emotion detection and sentiment analysis.

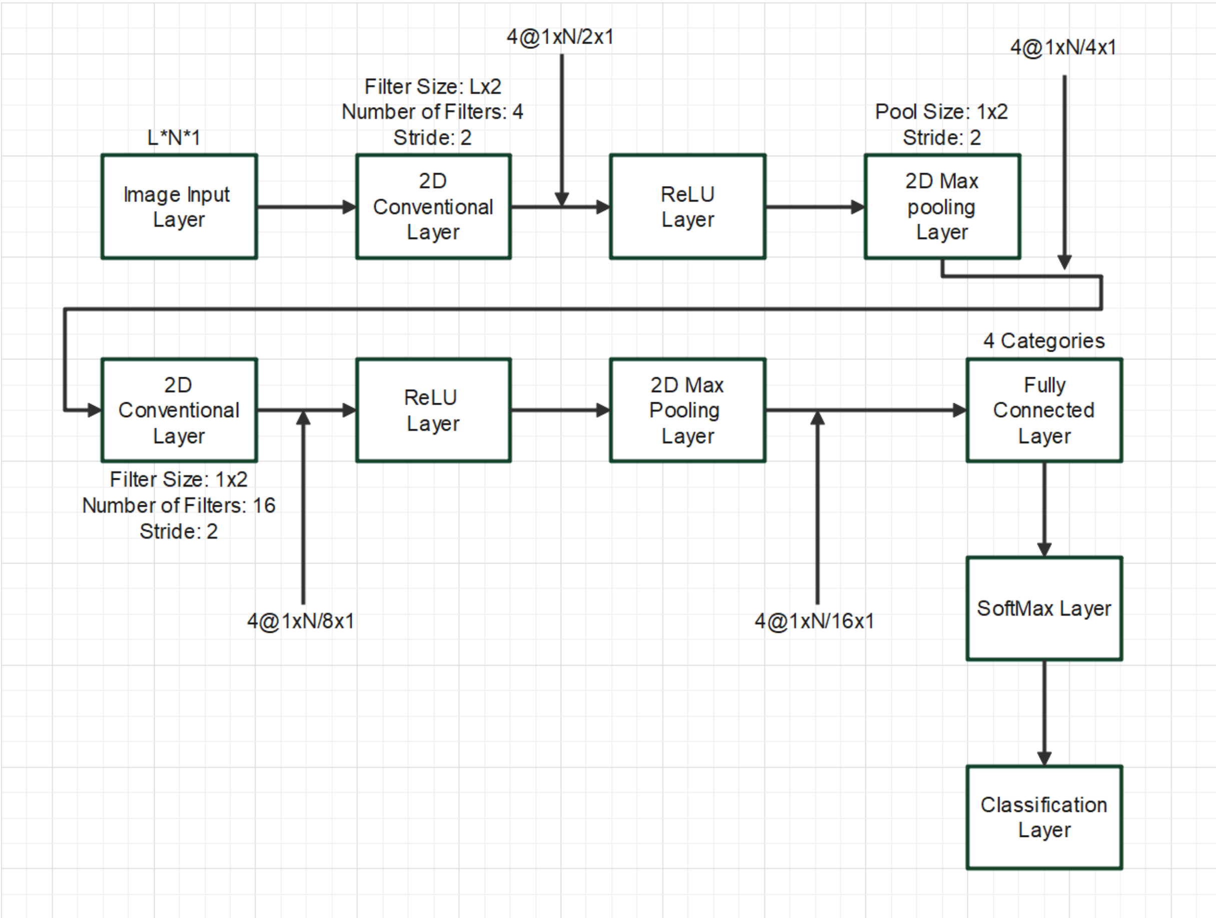

Figure 1 – CNN Flow Graph

1.3.1 What is Emotion Detection and Sentiment Analysis?

Emotion detection and sentiment analysis are closely related tasks that aim to understand the emotional tone and opinion conveyed in a piece of text.

- **Emotion Detection:** Identifies the specific emotions (e.g., joy, sadness, anger) expressed in text.
- **Sentiment Analysis:** Determines the overall sentiment (positive, negative, or neutral) of a text piece.

These tasks have numerous applications, including:

- Analyzing customer reviews
- Monitoring social media sentiment
- Improving chatbots and virtual assistants

1.3.2 How CNNs Work for Emotion Detection and Sentiment Analysis

Unlike the movie recommendation system using Count Vectorization and Cosine Similarity, CNNs offer a more advanced approach for emotion detection and sentiment analysis. Here's a breakdown:

1. **Text Preprocessing:** Similar to the movie recommendation system, textual data (e.g., tweets, reviews) needs preprocessing. This may involve cleaning, tokenization (splitting into words), and potentially embedding words into numerical vectors.
2. **Convolutional Layers:** CNNs are known for their ability to extract features from spatial data like images. However, for text, we can view sentences as a sequence of "words" where each word represents a feature.
   - The convolutional layers in a CNN for text take these word embeddings as input and learn to identify patterns within sequences. These patterns might capture emotional cues like exclamation marks or specific word combinations.
3. **Pooling Layers:** Pooling layers reduce the dimensionality of the data extracted by the convolutional layers. This helps to control overfitting and improve generalization.
4. **Fully-Connected Layers:** After feature extraction, fully-connected layers process the information from the convolutional layers. Depending on the task (emotion detection or sentiment analysis), there might be multiple output neurons with activations representing different emotions or sentiment classes (positive, negative, neutral).

1.3.3 Advantages of CNNs for Emotion Detection and Sentiment Analysis

- **Automatic Feature Extraction:** CNNs can automatically learn relevant features from the text data, eliminating the need for hand-crafted features like Count Vectorization.
- **Capability to Capture Long-Range Dependencies:** CNNs have the ability to capture relationships between words that might be further apart in a sentence, potentially leading to better understanding of context and emotional cues.

1.3.4 Conclusion

This section provided a brief introduction to how CNNs can be used for emotion detection and sentiment analysis. Compared to the simpler approach used in the movie recommendation system, CNNs offer a more powerful and flexible approach to analyze textual data and extract emotional information.

Note: This explanation serves as a high-level overview. The specific architecture and implementation details of CNNs for emotion detection and sentiment analysis can vary depending on the complexity of the task and the chosen model.

1.4 PROBLEM STATEMENT

Challenge: Accurately identify the emotions and overall sentiment expressed in textual data.

Context: The ever-growing volume of user-generated text data (e.g., social media posts, online reviews, customer feedback) presents a valuable opportunity to understand user opinions, attitudes, and emotions. However, manually analyzing this data is time-consuming and impractical.

Objective: Develop a robust and automated system to analyze textual data and extract emotional information. This system should be able to:

- **Classify emotions:** Accurately categorize the emotions expressed in text (e.g., joy, sadness, anger, fear, disgust).
- **Determine sentiment:** Identify the overall sentiment of the text (positive, negative, or neutral).

Current Limitations: Existing approaches based on traditional machine learning techniques like Count Vectorization and Cosine Similarity may not be suitable for capturing the complexities of human language and emotions.

Proposed Solution: Utilize Convolutional Neural Networks (CNNs) to analyze textual data. CNNs excel at extracting features from sequences, making them well-suited for processing text data and learning relevant patterns that indicate emotions and sentiment.

Potential Benefits:

- **Improved Customer Service:** Analyze customer reviews and feedback to understand customer satisfaction and identify areas for improvement.
- **Enhanced Social Media Monitoring:** Track public sentiment towards brands, products, or events on social media platforms.
- **Personalized User Experiences:** Recommend products or content based on the user's emotional state and preferences.
- **Advanced Chatbots and Virtual Assistants:** Develop chatbots that understand and respond to user emotions, leading to more engaging interactions.

Research Questions:

- What type of CNN architecture is most effective for emotion detection and sentiment analysis?
- What pre-processing techniques are necessary to optimize text data for CNN models?
- How can we handle the challenge of limited labeled data for emotion classification tasks?
- How can we combine CNNs with other techniques (e.g., attention mechanisms) to further improve performance?

Evaluation Metrics:

- Accuracy: Percentage of correctly classified emotions or sentiments.
- Precision: Proportion of positive predictions that are truly positive.
- Recall: Proportion of actual positive cases that are correctly identified.
- F1-score: Harmonic mean of precision and recall.

By addressing these challenges and research questions, we can develop a more robust and accurate CNN-based system for emotion detection and sentiment analysis, enabling deeper insights into human emotions expressed in textual data.

1.5 OBJECTIVE

The primary objective of using Convolutional Neural Networks (CNNs) for emotion detection and sentiment analysis is to develop a system that can automatically and accurately extract emotional information from textual data. This translates to achieving the following specific objectives:

1. **Accurate Emotion Classification:** Develop a CNN model capable of accurately classifying the emotions expressed in text data into predefined categories (e.g., joy, sadness, anger, fear, disgust). This objective focuses on assigning the correct emotional label to a piece of text.
2. **Reliable Sentiment Analysis:** Train a CNN model to effectively determine the overall sentiment (positive, negative, or neutral) conveyed in a text passage. This objective aims to identify the broader feeling or opinion expressed in the text.
3. **Feature Learning from Text:** Leverage the power of CNNs to automatically extract relevant features from textual data that are indicative of emotions and sentiment. Unlike traditional methods requiring manual feature engineering, CNNs can learn these features directly from the training data.
4. **Generalizability and Robustness:** Train a CNN model that performs well not only on the training data but also on unseen data. This objective emphasizes the model's ability to adapt and make accurate predictions on new text it hasn't encountered before.
5. **Scalability and Efficiency:** Design a CNN architecture that efficiently processes large volumes of textual data while maintaining a high level of accuracy. This objective ensures the system can handle the ever-growing amount of user-generated text data efficiently.

By achieving these objectives, a CNN-based system can provide valuable insights into user opinions, attitudes, and emotions expressed in textual data. This information can be used for various applications such as improving customer service, monitoring social media sentiment, and personalizing user experiences.

CHAPTER - 2

LITERATURE REVIEW

J.Islam et al[7] This paper presents a pioneering framework for visual sentiment analysis, leveraging transfer learning techniques to bridge this gap. Our approach entails initializing network models with hyper-parameters derived from a deep convolutional neural network, strategically employed to counteract overfitting tendencies. Through extensive experimentation utilizing a Twitter image dataset, we empirically demonstrate the superior performance of our model, surpassing prevailing state-of-the-art methodologies. Our findings not only contribute to advancing the field of online multimedia big data research but also furnish a robust methodology for decoding sentiment from visual content, particularly within the context of social media platforms. Moreover, we discuss the implications of our results in enhancing user experience, facilitating targeted content delivery, and understanding socio-emotional dynamics in online communities. This research lays a solid foundation for further exploration and development in the realm of visual sentiment analysis, offering promising avenues for future research and applications in sentiment-driven computational systems.
A. Severyn et al[8] present a cutting-edge deep learning system tailored for sentiment analysis of tweets. The primary innovation lies in introducing a novel model for initializing the parameter weights of the convolutional neural network, a critical aspect for training a precise model while circumventing the necessity of incorporating additional features. The approach involves utilizing an unsupervised neural language model to train initial word embeddings, which are subsequently fine-tuned by the deep learning model on a distant supervised corpus. In the final stage, the pre-trained parameters of the network are employed to initialize the model, which is then trained on the supervised training data provided by the official system evaluation campaign on Twitter Sentiment Analysis conducted by Semeval-2015. Evaluation against systems participating in the challenge on the official test sets reveals promising results, positioning our model among the top performers in both the phrase-level subtask A and the message-level subtask B. This empirical evidence underscores the practical efficacy of our solution, suggesting its potential applicability in real-world sentiment analysis tasks.

Q. You et al[9] address the significance of sentiment analysis in understanding online user-generated content within the realm of social media analytics. While previous research predominantly relied on textual sentiment analysis for tasks such as predicting political elections and measuring economic indicators, the emergence of additional images and videos as mediums for expressing opinions and experiences necessitates a comprehensive approach. Motivated by the imperative to harness large-scale social multimedia content for sentiment analysis, the authors propose a novel methodology that integrates state-of-the-art visual and textual sentiment analysis

techniques for joint visual-textual sentiment analysis. The approach involves fine-tuning a convolutional neural network (CNN) for image sentiment analysis alongside training a paragraph vector model for textual sentiment analysis. Extensive experimentation is conducted on both machine weakly labeled and manually labeled image tweets. The results demonstrate that leveraging joint visual-textual features yields superior performance compared to standalone textual and visual sentiment analysis algorithms, underscoring the efficacy of the proposed approach in capturing nuanced user sentiments towards events or topics within social media discourse. This research contributes significantly to the advancement of sentiment analysis methodologies tailored for the analysis of diverse and expansive social multimedia content.
In the landscape of Natural Language Processing (NLP), sentiment analysis stands as a prominent research area; however, its exploration within the Chinese language domain remains relatively scarce. This scarcity is primarily attributed to the absence of large, labeled corpora and robust models. To address this deficiency, C. Li et al[10] undertake the development of a Chinese Sentiment Treebank utilizing social data, comprising 13,550 labeled sentences sourced from movie reviews. Additionally, they propose a groundbreaking Recursive Neural Deep Model (RNDM) tailored for predicting sentiment labels through recursive deep learning. Focused on the task of classifying individual sentences based on overall sentiment, the model aims to discern whether a review exhibits positive or negative sentiment. Through extensive experimentation, particularly in predicting sentiment labels at the sentence level, the proposed RNDM model showcases significant performance improvements over commonly employed baselines such as Na¨ive Bayes, Maximum Entropy, and SVM. This research not only contributes to the enrichment of sentiment analysis methodologies

within the Chinese language domain but also underscores the efficacy of recursive deep learning techniques in capturing nuanced sentiment nuances in textual data.

C. **Baecchi et al[11]** delve into the realm of multimodal feature learning to tackle sentiment analysis of micro-blogging content, particularly focusing on short messages from platforms like Twitter that may include both text and images. The research is underpinned by recent advancements in two key areas: i) the development of language models based on neural networks, exemplified by techniques such as Skip-gram and Denoising Autoencoders, renowned for their efficacy in processing web-scale text corpora and capturing syntactic and semantic word similarities; and ii) unsupervised learning of robust visual features using neural networks, capable of extracting meaningful features from partially observed images, accommodating occlusions, noise, and heavy modifications. To leverage these advances, the authors propose a novel architecture integrating these neural networks, which is subsequently tested on various standard Twitter datasets. The experimental results validate the efficiency and effectiveness of the proposed approach, demonstrating its capability to achieve favorable classification results in the realm of micro-blogging sentiment analysis. This research contributes to the advancement of multimodal sentiment analysis methodologies, offering promising avenues for further exploration in leveraging both textual and visual information for sentiment analysis tasks in micro-blogging platforms.

In the paper, **Yang et al[12]** investigate the burgeoning field of visual sentiment analysis, which has garnered increasing attention with the growing tendency to convey emotions through images. While existing methodologies typically assign a single dominant emotion to each image, the authors confront the challenge of sentiment ambiguity by introducing label distribution learning (LDL), recognizing that images often evoke multiple emotions. They propose two novel algorithms built upon conditional probability neural networks (CPNN). Initially, they present the Binary CPNN (BCPNN), which encodes image labels into binary representations, enhancing input for the neural network. Subsequently, they introduce the Augmented CPNN (ACPNN), which incorporates label noise and augments affective distributions during training. Recognizing the scarcity of large-scale datasets annotated with multiple affective labels, the authors construct two new datasets, named Flickr LDL and Twitter LDL, featuring over 20,745 images with multiple affective labels, surpassing existing datasets in scale. Experimental results demonstrate the superiority of the proposed methods over state-of-the-art approaches, showcasing their efficacy in predicting multi-emotions evoked in images. Leveraging mid-level and deep features on various datasets, including Abstract Paintings, Emotion6, Flickr LDL, and Twitter LDL, the authors establish that ACPNN outperforms BCPNN, with both models exhibiting superior performance compared to CPNN and other contrastive methods. This research contributes to advancing the field of visual sentiment analysis by addressing the inherent ambiguity in image sentiment and introducing effective methodologies for leveraging label distribution learning in sentiment analysis task.

The authors **Vadicamo et al[13].** conducted pioneering research in the domain of sentiment analysis, specifically focusing on the integration of textual and visual modalities. Departing from the conventional reliance on textual data alone, they embarked on a novel exploration of predicting sentiments from multimedia content. This research endeavor tackled the formidable challenge of training a visual sentiment classifier using an extensive corpus of user-generated, unlabeled content, comprising a staggering dataset of over 3 million tweets embedding both textual and visual elements. Employing a cross-media learning paradigm, the authors ingeniously leveraged the sentiment polarity inherent in textual content to facilitate the training of a deep Convolutional Neural Network (CNN) tailored for predicting sentiment polarity in associated images. This methodology marks a seminal departure in sentiment analysis research, representing the first concerted effort to bridge the textual-visual gap in sentiment prediction tasks.The empirical evaluation conducted by the authors not only elucidated the viability of their approach but also underscored its efficacy in accommodating the inherent noise and weak correlation often encountered between textual and visual modalities. Through meticulous experimentation, they demonstrated the robustness of their model in accurately predicting sentiment polarity in images under uncontrolled conditions, colloquially referred to as in the wild. A notable contribution of this research lies in the authors' provision of their comprehensive dataset and trained models to the research community, fostering collaborative endeavors and enabling further exploration in the burgeoning field of visual sentiment analysis.

In their research, the authors **Siersdorfer et al[14]** delve into the intricate relationship between the sentiment conveyed through metadata and the visual content of images within the social photo sharing platform Flickr. They adopt a multi-faceted approach by analyzing both the bag-of-visual words representation and the color distribution of images, while simultaneously tapping into the SentiWordNet thesaurus to derive numerical sentiment values from accompanying textual metadata. Employing discriminative feature analysis grounded in information theoretic principles, the authors rigorously investigate the connection between visual features and sentiment. Leveraging machine learning techniques, they undertake the challenging task of predicting the sentiment polarity of images. Through an extensive empirical study encompassing a vast dataset comprising over half a million Flickr images, the authors unveil a substantial correlation between sentiment and visual features. Their findings offer promising insights into the feasibility of estimating sentiment polarity within images, thereby advancing the understanding of sentiment analysis in visual content.

Peng et al[15] delve into two novel aspects concerning photos and human emotions. Firstly, they conduct psychovisual studies revealing that different individuals exhibit varied emotional responses to the same image, a departure from previous research that typically focused on predicting a single dominant emotion per image. Furthermore, their studies demonstrate that a single individual may experience multiple emotional reactions to a single image. This shift towards predicting emotions in distributions rather than singular dominant emotions holds significance for numerous applications. Secondly, the authors demonstrate the potential to manipulate the evoked emotion of an image by adjusting color tone and texture-related features, thereby influencing the direction of emotional change. They introduce the Emotion6 database, which contains distributions of emotions, paving the way for a new paradigm where emotional responses to images are represented as distributions rather than single emotions. The authors propose methods for estimating emotion distributions for images and present a technique for modifying images to align their emotional distribution with a target image. Their novel emotion predictor, CNNR, outperforms existing methods, including Support Vector Regression (SVR) with features from prior works and the optimal dominant emotion baseline.Lastly, the authors introduce the Emotion6 database, providing ground truth data on valence, arousal, and probability distributions of evoked emotions. This comprehensive dataset facilitates further research into the intricate relationship between visual stimuli and human emotional responses, contributing to advancements in emotion recognition and understanding

In this report, the authors **Mikels et al [16]** provide categorical data that enhance the utility of the International Affective Picture System (IAPS) for studying emotions from a discrete categorical perspective. Consistent with prior findings, minimal gender differences were observed in the emotional categorization of IAPS images. The data reveal that a significant portion of the images elicit either single discrete emotions or emotions representing blends of discrete emotions. Importantly, these data enable researchers to explore the behavioral effects of discrete single emotions versus blended emotions. Various studies have demonstrated that different emotions have distinct cognitive and behavioral consequences. For instance, sadness and anger have differing effects on judgments of situational-caused versus human-caused events, while fear and anger influence perceptions of future events differently. Fear tends to lead to risk aversion, whereas anger leads to risk-seeking behavior. Consequently, understanding the effects of blended emotions versus single discrete emotions on cognition and behavior becomes crucial. Although the conclusions drawn are based on the authors' method of analysis, alternative methods could be employed to interrogate the data, especially concerning the intensity of discrete emotional labels. The data set allows researchers flexibility in employing alternative methods tailored to their specific research inquiries. Interestingly, the data suggest fewer instances of single positive discrete emotions elicited by the images, which may align with theoretical perspectives positing that positive emotions are fewer, more diffuse, and subtler. Moreover, the absence of images eliciting anger alone may stem from methodological challenges associated with anger induction through passive viewing of static images. The authors emphasize the importance of considering diverse emotion induction techniques, as the effectiveness of a technique varies depending on the specific emotion targeted. While the IAPS has traditionally been utilized for studying dimensional aspects of emotions, this report introduces new categorical data that augment the utility of the image set for research paper.

In this paper, **K Patel et al[17]** conduct a meticulous and systematic survey aimed at analyzing the current state-of-the-art methodologies for facial emotion recognition (FER) in static images, along with the various parameters that influence the efficacy of these approaches. They introduce a taxonomy based on different techniques utilized for face detection, feature extraction, and emotion classification. Moreover, the authors review the utilization of various facial expression databases in FER studies.Acknowledging the significance of machine and deep learning algorithms in advancing computer vision applications, particularly in facial sentiment analysis, the authors identify technological challenges faced by deep learning-based FER models, such as under-fitting or over-fitting due to insufficient training data or expression diversity. Motivated by these challenges, the paper presents a comprehensive survey of state-of-the-art Artificial Intelligence techniques, encompassing datasets and algorithms aimed at addressing these issues. Furthermore, the paper provides a taxonomy of existing facial sentiment analysis strategies and reviews novel machine and deep learning networks specifically tailored for FER based on static images. The authors critically evaluate the merits and demerits of these approaches, summarizing their methodologies. Finally, the paper outlines open issues and research challenges in the field of FER, including the exploration of FER in 3D face shape models and recognizing emotions in images under occlusion. Real-time FER remains a formidable task, and the authors underscore the necessity for continued research efforts. They express their intention to further investigate FER in videos using more advanced deep learning techniques in the future.

In this chapter, the authors **Alex D. Torres et al[18]**conducts sentiment analysis of medical patients through facial emotion recognition, utilizing machine learning-based image classifiers. The system design comprises three primary components: a video processor, a face detector, and an image classifier. The video processor captures frames from a video feed, converts them to grayscale images, and passes them to the face detector, which performs face detection by cropping the image around the face. The cropped face images are then resized and fed into the image classifier for facial emotion recognition.The author employs a convolutional neural network (CNN) model, extending the state-of-the-art model by utilizing pre-processed video stream frames as input to the CNN. Additionally, the author introduces an optimization step during training to combat saturation in validation accuracy, dynamically adjusting the learning rate to converge on optimal trainable parameters.The implemented system achieves inference accuracy comparable to state-of-the-art facial emotion classifiers and is deployed on a remote, cloud-based GPU platform, leveraging CUDA programming techniques for accelerated model training. This enables real-time operation of the system. Furthermore, the author conducts a case study using the proposed facial emotion recognition system to automatically assess patients' pain intensity levels. The framework, named "DeepPain," correlates pain intensity with probabilities of facial emotion categories. Specifically, probabilities of "fearful" and "sad" emotions are identified as closely related to pain intensity and utilized as inputs to the DeepPain framework. The DeepPain framework facilitates fine-grained detection and understanding of pain intensity levels, which holds significant implications for healthcare applications.

In this paper, **Dinesh Kalla et al[19]**develops a deep convolutional neural network (CNN) model for facial emotion recognition, trained on a substantial dataset comprising 36,000 grayscale images of faces annotated with seven emotion categories: happy, sad, angry, surprised, fearful, disgusted, and neutral. The methodology involves the utilization of four convolutional layers followed by two fully connected layers to extract hierarchical visual features, which are subsequently classified into probabilities across the emotion categories.The analysis indicates that the model achieves a reasonable validation accuracy of 65% in categorizing basic emotions from subtle facial expressions. However, performance significantly deteriorates when tested on more diverse real-world data reflecting variations in poses, illumination, resolution, and background conditions.To address these challenges, the paper discusses various opportunities for improvement, including techniques such as data augmentation to enhance diversity, personalized adaptation to accommodate individual facial variations, ensemble methods for robustness, three-dimensional modeling for richer representation, graph networks for contextual understanding, and self-supervised representation learning for feature extraction. Facial emotion recognition holds immense promise for applications in human-computer interaction across diverse domains such as healthcare, education, and transportation. By addressing the challenges highlighted and leveraging the discussed techniques, advancements in this field could lead to more effective and versatile systems with broader real-world applicability.

In this thesis, the authors **Md. Forhad Ali et al[20]** explores the capability of machines to recognize human emotions, aiming to enable them to act more like humans and potentially prevent undesirable occurrences. The author emphasizes the increasing intelligence of machines and their capacity to mimic human sentiment. To achieve this, a comprehensive framework is developed, informed by extensive research and fieldwork, to determine emotion expression patterns. The framework utilizes deep learning Convolutional Neural Network (CNN) algorithms implemented with tools such as Keras and TensorFlow for emotion recognition in real images. Decision tree techniques are also introduced to visually delineate the results and procedures, aiding in determining the predominant emotions and their respective probabilities. By employing these techniques, the study demonstrates the potential for machines to accurately identify emotions in images, enabling them to react appropriately and potentially prevent unwanted occurrences. This advancement suggests that machines could serve as effective replacements for humans in certain contexts, given their ability to recognize and respond to human emotions with increasing accuracy.

In this research, the authors **Wikarsa et al[21]** addresses the need for an application capable of detecting user emotions on Twitter, leveraging text mining techniques. Despite previous studies classifying sentiments into three categories (positive, negative, neutral), this research endeavors to classify emotions into six categories: happiness, sadness, anger, disgust, fear, and surprise. The text mining application follows three main phases: preprocessing, processing, and validation. In the preprocessing phase, activities include case folding, cleansing, stop-word removal, emoticons conversion, negation conversion, and tokenization applied to both training and test data. Morphological analysis is performed to build multiple models for sentiment analysis. The processing phase involves weighting and classification utilizing the Naive Bayes algorithm on the validated model. Finally, in the validation phase, the application measures accuracy using 10-fold cross-validation. The findings reveal that the application achieves an accuracy of 83 percent for 105 tweets. However, to enhance accuracy further, the author suggests the necessity for a better model in the training data. This research represents a significant step towards the development of an effective tool for detecting user emotions on Twitter, offering potential applications in sentiment analysis and social media analytics.

In their research paper, **Perikos et al. [22]** delve into the intricate domain of recognizing emotions in textual data through the lens of ensemble learning. With a focus on enhancing machine comprehension of human emotions, the authors strive to imbue machines with a deeper understanding of human-like sentiment, thereby potentially mitigating adverse occurrences. Emphasizing the escalating intelligence of machines and their ability to mimic human sentiment, the authors introduce a comprehensive framework informed by meticulous research and empirical investigations. This framework intricately employs ensemble learning techniques, harnessing the collective strengths of multiple classifiers. Through the amalgamation of diverse classifiers such as support vector machines (SVM), decision trees, naive Bayes, and neural networks, the authors aim to decode nuanced emotional cues embedded within textual content. Moreover, the authors introduce adept feature engineering methodologies tailored to capture subtle emotional nuances present in textual data. Through rigorous experimentation on benchmark emotion datasets, the study unveils the efficacy of the ensemble approach, outperforming individual classifiers consistently. This research not only advances the frontier of emotion recognition in text but also signifies a significant stride towards empowering machines to recognize and respond to human emotions adeptly. Consequently, the findings hint at the potential for machines to serve as competent surrogates for human intervention, especially in contexts where discerning and reacting to emotions in text are paramount.

In their study, **LeCompte et al.[23]** embark on the exploration of sentiment analysis in the realm of social media discourse, specifically focusing on tweets enriched with emoji data. With an overarching goal of unraveling the sentiment embedded within these succinct yet expressive messages, the authors aim to enrich machine

CHAPTER - 3

PROJECT PROCEDURE

1. Data Collection and Preprocessing

- **Facial Expression Data:**
  - Identify and select appropriate facial expression datasets based on your project goals. Consider factors like dataset size, emotion categories covered, image quality, and licensing restrictions. Some popular options include:
    - JAFFE: Provides grayscale images of Japanese female faces with seven emotions (neutral, happy, sadness, surprise, anger, disgust, fear). Limited size and diversity.
    -

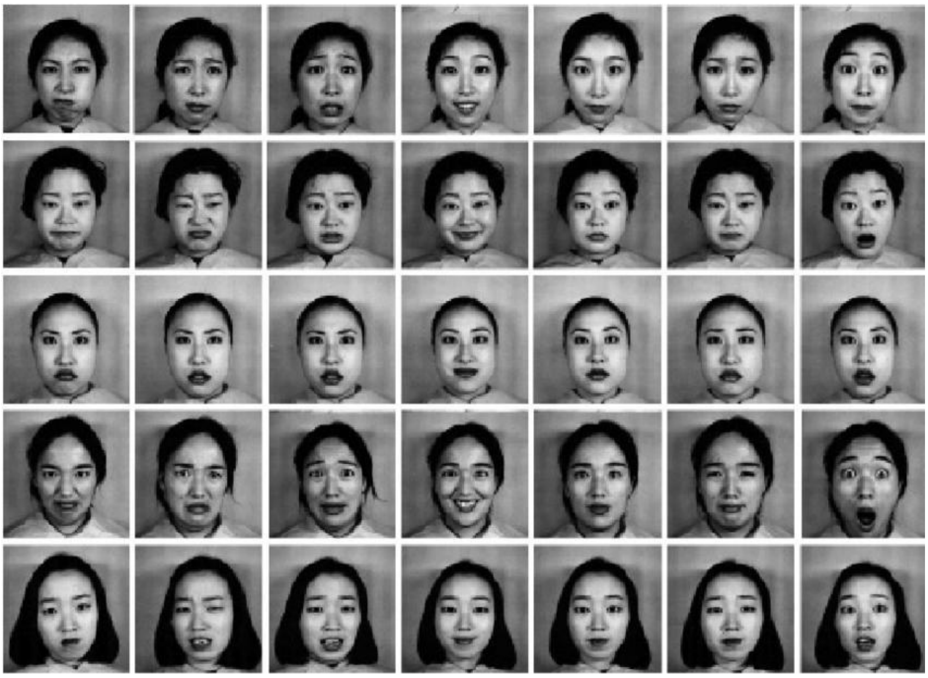

Figure 3(a) - JAFFE (example image)

    - CK+: Offers grayscale images from various subjects with posed and non-posed expressions, covering six basic emotions (happiness, sadness, surprise, anger, fear, disgust). More diverse than JAFFE but still limited in size.
    - FER-2013: A large-scale dataset with grayscale facial images sourced from Kaggle, labeled with seven emotions (happy, sad, angry, surprised,

fearful, disgusted, neutral). Ideal for training deep learning models due to its size.

- **Text Data:**
    - Choose textual datasets that align with your sentiment analysis focus. Options include:
        - IMDB movie reviews: Labeled with positive or negative sentiment for movie reviews.
        - Election Tweets: Tweets related to political elections, candidates, and voter opinions. Can be used to analyze sentiment towards specific topics or candidates.
        - General Tweets: Covers a broad spectrum of user-generated content, opinions, and interactions. Useful for general sentiment analysis tasks.
- **Preprocessing Techniques:**
    - **Facial Images:**
        - Convert images to grayscale for consistency.
        - Resize images to a standard size to ensure compatibility with your models.
        - Normalize pixel intensity values (e.g., between 0 and 1) for improved training efficiency.
        - Consider techniques like data augmentation (e.g., random flipping, cropping, rotation) to artificially increase dataset size and improve model robustness to variations.
    - **Text Data:**
        - Remove punctuation marks, stop words (common words with little meaning), and special characters for cleaner text analysis.
        - Convert all text to lowercase for case-insensitivity.
        - Tokenize text into individual words or phrases for further processing.
        - Explore techniques like stemming (reducing words to their root form) or lemmatization (converting words to their dictionary form) depending on your chosen model and desired level of detail.

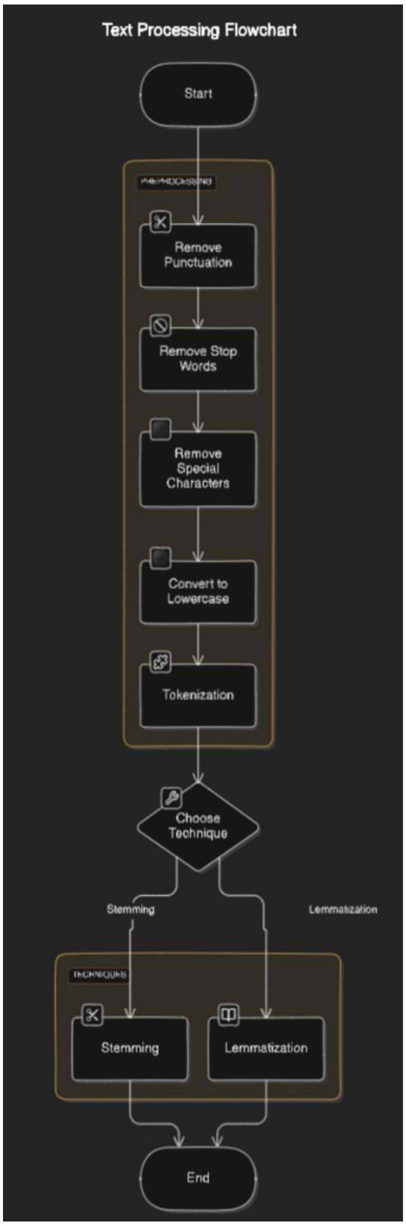

Figure 3(b) - Text Processing

2. Model Selection and Training

- **Facial Expression Recognition:**
  - Convolutional Neural Networks (CNNs) are a strong choice due to their ability to automatically learn spatial features from images. Popular CNN architectures for image classification include:
    - VGG16: A deep convolutional neural network with good performance on various image classification tasks.
    - ResNet: A family of CNN architectures known for addressing the vanishing gradient problem in deep networks.
    - EfficientNet: A family of lightweight and efficient CNN models offering good accuracy with fewer parameters.
- **Sentiment Analysis:**

Figure 3(c) - RNN

  - Recurrent Neural Networks (RNNs) or transformers (like BERT) are well-suited for analyzing sequential data like text. Consider these options:
    - Long Short-Term Memory (LSTM): A type of RNN capable of learning long-term dependencies in text sequences.
    - Gated Recurrent Unit (GRU): Another RNN variant with a simpler architecture compared to LSTMs, potentially faster to train.
    - Bidirectional RNNs: Process text in both forward and backward directions to capture context from both sides of a word.
    - BERT (Bidirectional Encoder Representations from Transformers): A pre-trained transformer model that can be fine-tuned for various NLP tasks, including sentiment analysis.

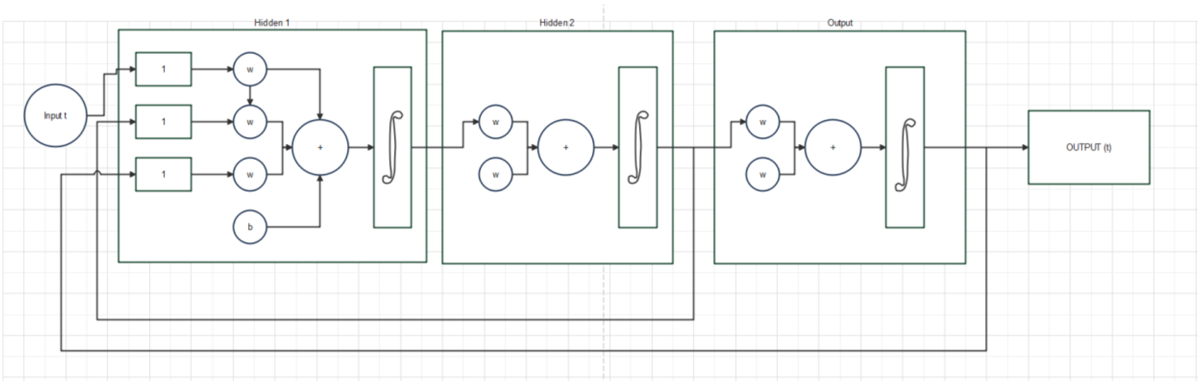

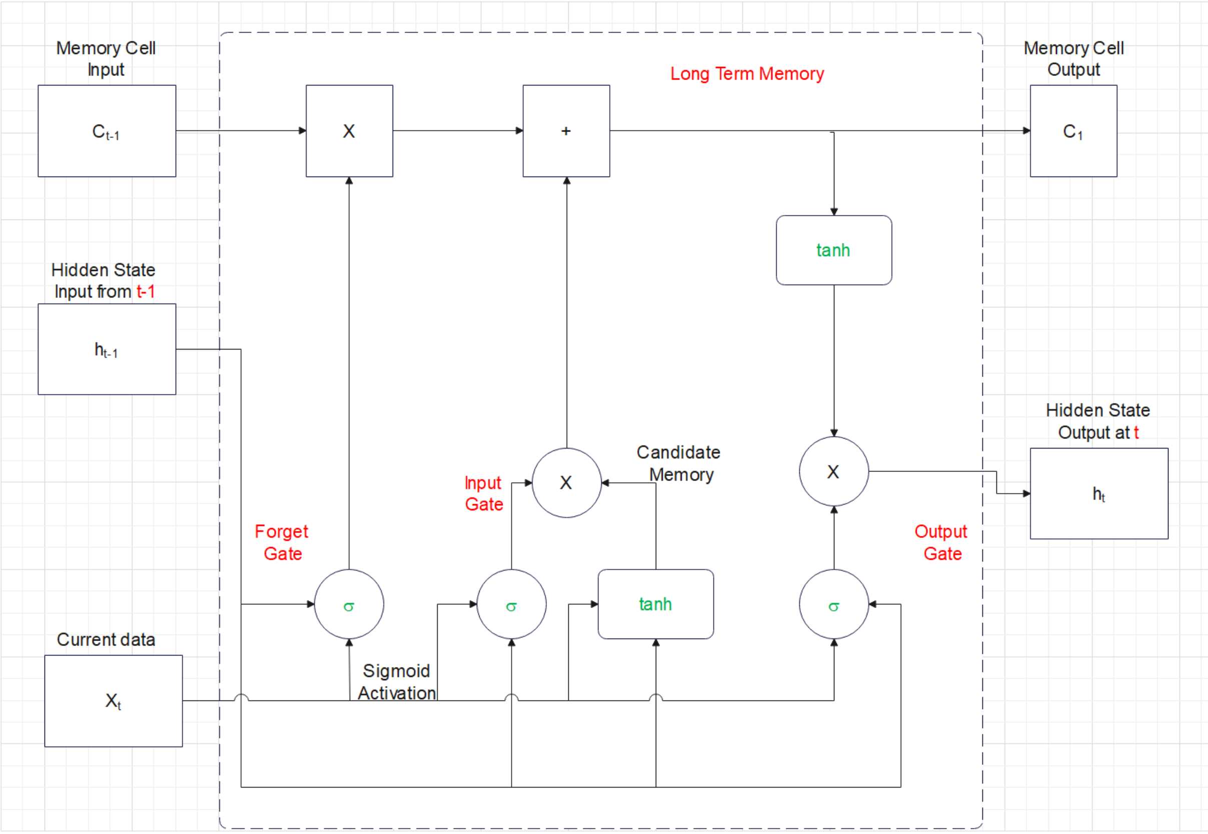

Figure 3(d) - LSTM

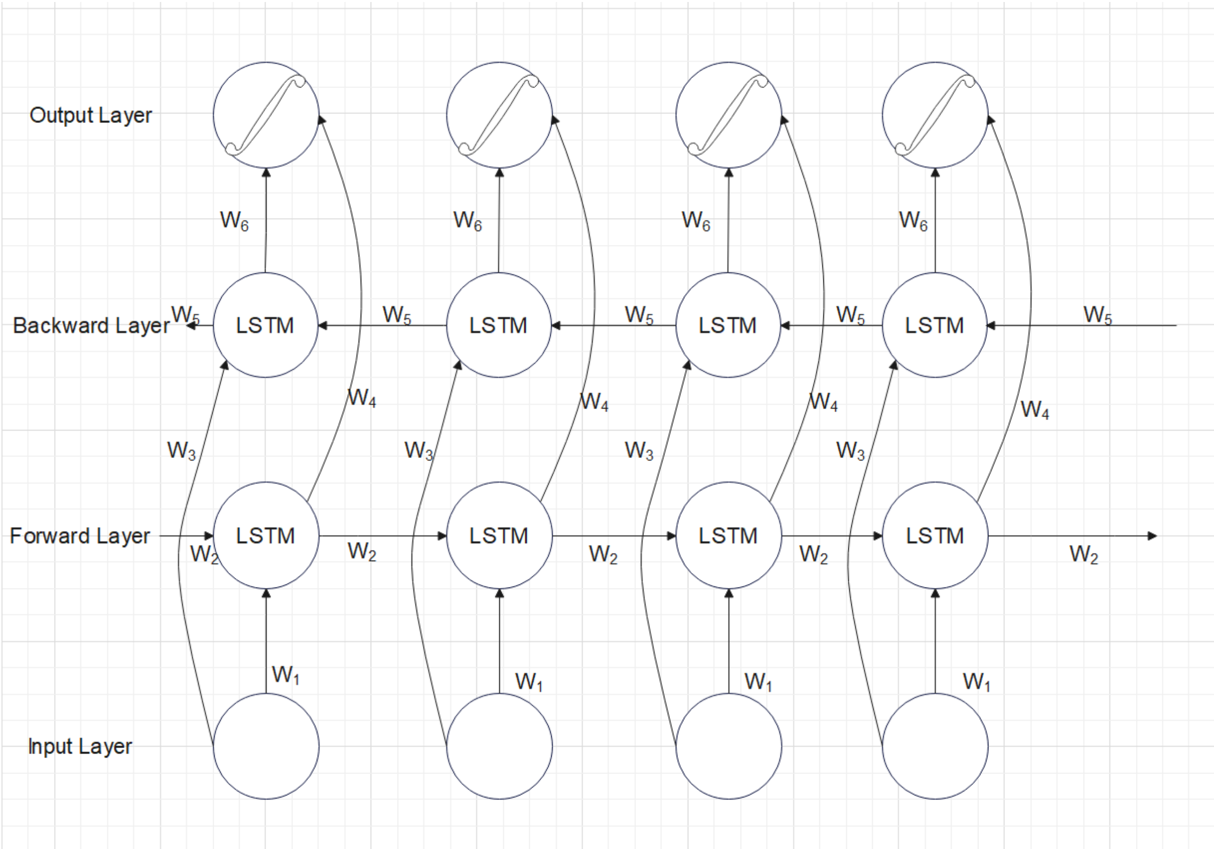

Figure 3(e) - Bi-LSTM

- **Training Process:**
    - Split the preprocessed data into training, validation, and testing sets. The training set is used to train the model, the validation set is used for hyperparameter tuning, and the testing set is used for final evaluation of the model's performance on unseen data. Common splits are 80% for training, 10% for validation, and 10% for testing.
    - Train the chosen model(s) on the training data. Use an optimizer (e.g., Adam) to adjust model weights during training and a loss function (e.g., categorical cross-entropy) to measure the difference between predicted and actual labels.
    - Monitor the model's performance on the validation set during training. Use techniques like early stopping to prevent overfitting (the model memorizing training data instead of learning general patterns).
    - Fine-tune hyperparameters (learning rate, number of epochs, etc.) based on validation set performance.

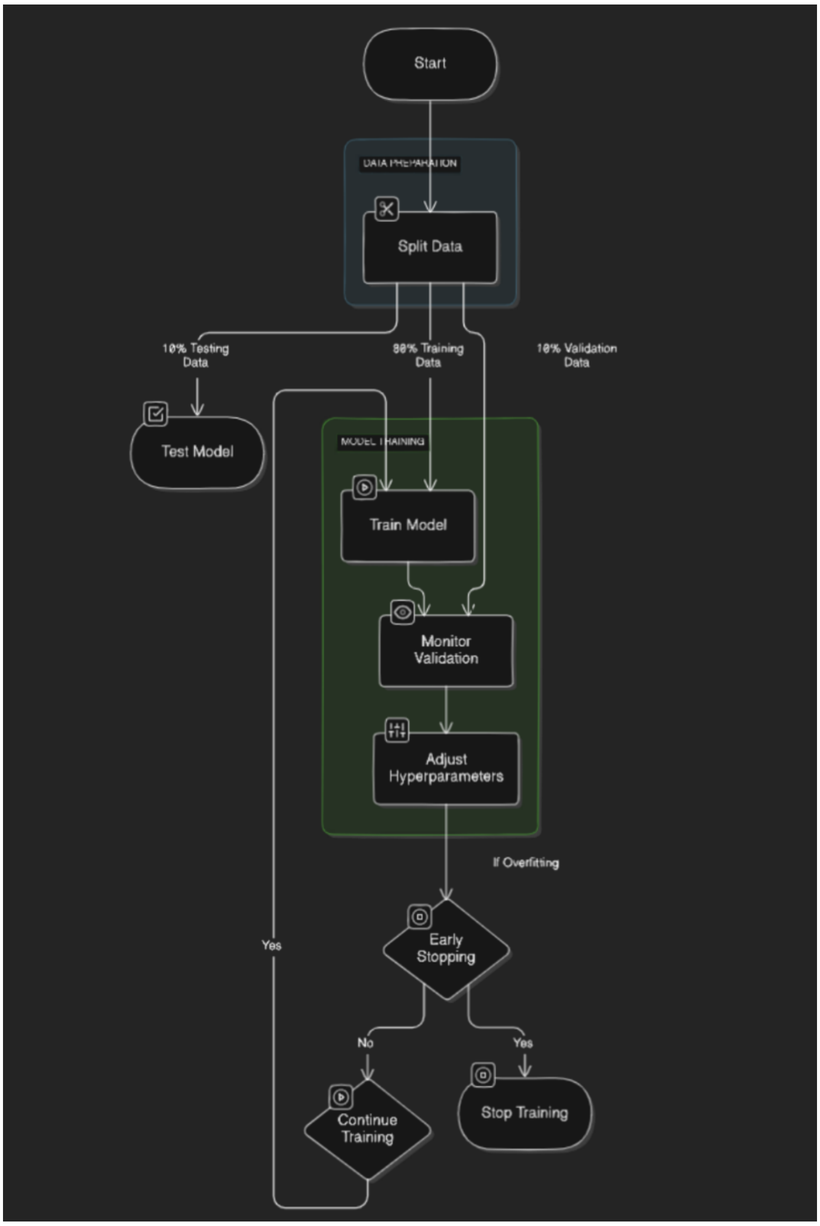

Figure 3(f) - Training Process

3. Evaluation and Analysis

- **Evaluation Metrics:**
    - Choose appropriate metrics to assess the performance of your models for both facial expression recognition and sentiment analysis. Here's a breakdown for each task:
        - **Facial Expression Recognition:**
            - Accuracy: Overall percentage of correctly classified facial expressions.

        - Precision: Measures how many predictions of a specific emotion were actually correct for that emotion.
        - Recall: Indicates how well the model identified all instances of a particular emotion.
        - F1-score: Harmonic mean of precision and recall, providing a balanced view of model performance.
    - **Sentiment Analysis:**
        - Accuracy: Proportion of correctly classified sentiments (positive, negative, or neutral, depending on your dataset).
        - Precision and Recall (as defined above) can be calculated for each sentiment class.
        - Confusion Matrix: Visualizes the model's performance by showing how many instances of each class were predicted as belonging to other classes.
- **Comparative Analysis:**
    - Train and evaluate multiple models (different CNN architectures or RNN variants) to compare their performance. This helps you identify the model that best suits your data and project goals.
- **Error Analysis:**
    - Analyze the types of facial expressions or text samples that the models struggle with the most. This can involve manually reviewing misclassified examples to understand potential reasons for errors.
    - Identify patterns in the errors: Are there specific emotions or words that are consistently misclassified? This analysis can guide further data augmentation or model improvement strategies.

4. Integration and Refinement

- **Model Integration:**
    - Once you have well-performing individual models for facial expression recognition and sentiment analysis, integrate them into a single system. The system should take an input (facial image and/or text data) and provide an output that combines the sentiment analysis result with the recognized emotion(s) from the facial expression.
- **Ensemble Learning (Optional):**
    - Consider techniques like ensemble learning to potentially improve the overall performance of your system. Ensemble methods combine predictions from multiple models to potentially achieve better accuracy than any single model.
- **Refinement Strategies:**
    - Based on the evaluation and analysis results, refine your system:
        - Try different model architectures or hyperparameter configurations.
        - Experiment with data augmentation techniques to create a more diverse dataset and potentially improve modelgeneralizability.
        - Explore transfer learning, where you leverage a pre-trained model on a large dataset (e.g., pre-trained CNN for facial features or pre-trained BERT for sentiment analysis) and fine-tune it on your specific data. This can be particularly beneficial when your own dataset is limited.

5. Testing and Deployment

- **Testing on Unseen Data:**
    - Evaluate the integrated system on unseen data to assess its ability to generalize to real-world scenarios beyond the training data. This helps ensure the system's robustness and applicability in practical situations.
- **Deployment (Optional):**
    - Depending on your project goals, you can consider deploying the system as a web application, mobile app, or API for wider use. This allows others to interact with your system and analyze sentiment or emotions based on facial expressions and text data.

6. Reporting and Documentation

- **Comprehensive Report:**
    - Prepare a detailed report summarizing your project, including:
        - Objectives and methodology used.
        - Data collection and preprocessing techniques.
        - Model selection, training, and evaluation processes.
        - Detailed results, discussion of findings, and limitations of your system.
        - Potential future work directions and applications of your project.
- **Code and Model Documentation:**
    - Document your code for reproducibility, allowing others to understand your implementation and potentially build upon your work.
    - Document the trained models, including details like architecture, hyperparameters, and performance metrics. This information can be valuable for future reference or sharing with collaborators.

By following these steps and conducting a thorough evaluation and analysis, you can develop and refine a robust system for sentiment analysis and emotion recognition using facial expressions and text data.

CHAPTER - 4

WORK DONE

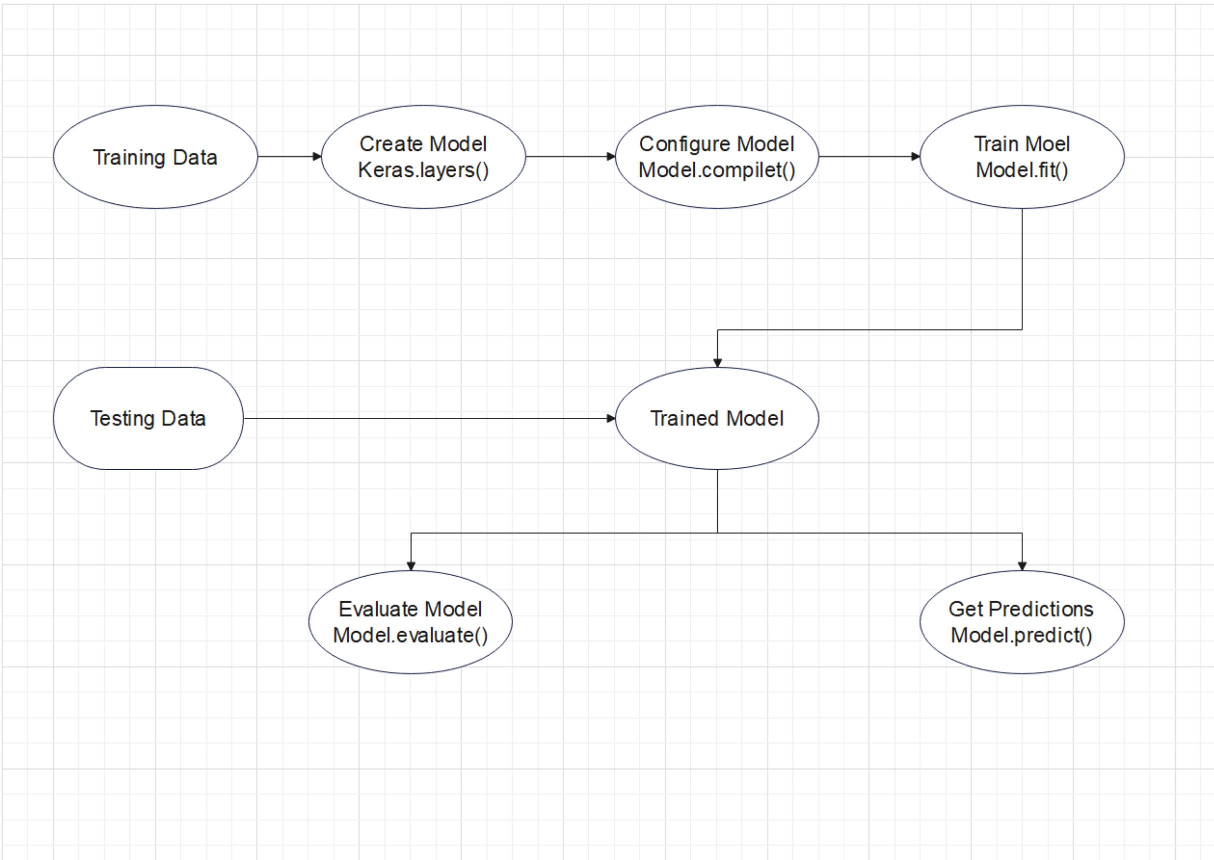

Figure 4(A) - Custom CNN Model

```
from keras.utils import to_categorical
from keras_preprocessing.image import load_img
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D
import os
import pandas as pd
import numpy as np
```

- **`from keras.utils import to_categorical`:** This import is used to convert categorical data (like genre labels) into a one-hot encoded format, which is often necessary for neural network training. However, since the reference talks about movie data that might not involve image classification, this import might not be directly applicable.

- **`from keras.preprocessing.image import load_img`:** This import is used to load images from the file system. It's relevant if you're working with image data for movie analysis (e.g., movie posters).
- **`from keras.models import Sequential`:** This import allows you to build a sequential neural network model, which is a series of layers stacked one after another.
- **`from keras.layers import Dense, Conv2D, Dropout, Flatten`:** These imports provide the building blocks for your neural network model:

  - **`Dense`:** Creates a fully connected layer with a specified number of neurons.
  - **`Conv2D`:** Creates a convolutional layer for processing 2D data like images (assuming image data).
  - **`Dropout`:** Introduces dropout regularization to prevent overfitting.
  - **`Flatten`:** Reshapes a multi-dimensional tensor (e.g., image data after convolutional layers) into a single dimension for feeding into fully connected layers.

```
from PIL import Image


TRAIN_DIR = 'images/train'
TEST_DIR = 'images/test'


def createdataframe(dir):
    image_paths = []
    labels = []
    for label in os.listdir(dir):
        for imagename in os.listdir(os.path.join(dir,label)):
            image_paths.append(os.path.join(dir,label,imagename))
            labels.append(label)
        print(label, "completed")
    return image_paths,labels
```

`from PIL import Image`: This line imports the `Image` class from the Python Imaging Library (PIL). PIL is a popular library for image processing tasks in Python. The `Image` class provides functionalities for loading, manipulating, and saving images.

- **`TRAIN_DIR = 'Images/train'`** •

  - This line assigns the string `'Images/train'` to a variable named `TRAIN_DIR`.
  - The single quote (`'`) is used to define a string literal.

- `TRAIN_DIR` likely refers to the directory containing your training images. The path indicates a folder named `train` within a parent folder named `Images`.

- **`TEST_DIR = 'images/test •'**

  - Similar to the first line, this assigns the string `'images/test'` to a variable named `TEST_DIR`.
  - The path here suggests a folder named `test` within a directory named `images` (lowercase `i`).

- The movie_id, title, synopsis, genres, keywords, cast, and crew columns were kept after the dataset was cleaned and processed to ensure their continued use in analysis. The "overview" column gives a brief synopsis of the film, "genres" identifies the genres connected to it, "keywords" contains pertinent search terms, and "cast" and "crew" list the actors and crew people who worked on it. These crucial elements of each movie will be the focus of a thorough examination of this condensed, improved dataset.

- **Initializing Empty Lists:**

  - `image_paths = []`: This creates an empty list named `image_paths` to store the full paths of all the images found in the directory.
  - `labels = []`: This creates an empty list named `labels` to store the corresponding labels for each image.

- **Iterating Through Subdirectories:**

  - `for label in os.listdir(dir):`: This loop iterates over all the subdirectories (folders containing images) within the directory specified by `dir`. For each subdirectory, the variable `label` will hold the name of that subdirectory.

- **Iterating Through Images:**

  - `for imagename in os.listdir(os.path.join(dir,label)):`: This nested loop iterates over all the files (assumed to be images) within the current subdirectory (`label`). The variable `imagename` will hold the name of each image file.

- **Building Image Paths:**

  - `image_paths.append(os.path.join(dir,label,imagename))`: This line constructs the full path for each image by joining the base directory (`dir`), the subdirectory name (`label`), and the image filename (`imagename`). The `os.path.join` function helps create platform-independent file paths. The constructed path is then appended to the `image_paths` list.

- **Assigning Labels:**

- `labels.append(label)`: Since the subdirectory name (`label`) is assumed to represent the image class or category, this line appends the `label` to the `labels` list. This creates a mapping between the image path and its corresponding label.

- **Printing Progress (Optional):**

  - `print(label, "completed")`: This line (assuming it's not commented out) prints a message indicating the completion of processing a subdirectory (label). This can be helpful for tracking progress, especially for large datasets.

- **Returning Results:**

  - `return image_paths,labels`: After iterating through all subdirectories and images, the function returns two values:
    - `image_paths`: The list containing all the full image paths.
    - `labels`: The list containing the corresponding labels for each image.

```
train = pd.DataFrame()
train['image'], train['label'] = createdataframe(TRAIN_DIR)
```

1. Importing pandas (assumed):

While not explicitly shown, this code likely assumes you have imported the `pandas` library using `import pandas as pd`. Pandas is a powerful library for data manipulation and analysis in Python.

2. Creating an Empty DataFrame:

This line creates an empty pandas DataFrame named `train`. A DataFrame is a two-dimensional tabular data structure with labeled columns.

3. Populating the DataFrame with Image Paths and Labels:

- It calls the `createdataframe` function, passing the value of the variable `TRAIN_DIR` (which presumably holds the path to your training image directory) as an argument.
- The function returns two lists: `image_paths` (containing full image paths) and `labels` (containing corresponding labels).
- These returned lists are unpacked and assigned to two new columns in the `train` DataFrame:

  - `train['image']`: This assigns the list of image paths to a new column named `'image'` in the DataFrame.
  - `train['label']`: This assigns the list of labels to a new column named `'label'` in the DataFrame.

```python
print(train)
```

```python
test = pd.DataFrame()
test['image'], test['label'] = createdataframe(TEST_DIR)
```

```python
print(test)
print(test['image'])
```

- **Printing Training Data (Optional):**
    - `print(train)`: Prints the contents of the `train` DataFrame for inspection.
- **Creating Testing DataFrame:**
    - `test = pd.DataFrame()`: Creates an empty DataFrame named `test` for testing data.
    - `test['image'], test['label'] = createdataframe(TEST_DIR)`: Similar to training, populates `test` DataFrame with data from the testing directory (`TEST_DIR`).
- **Printing Testing Data (Optional):**
    - `print(test)`: Prints the contents of the `test` DataFrame.
    - `print(test['image'])`: Optionally prints only the image paths from the testing data.

```python
from tqdm.notebook import tqdm
```

```python
def extract_features(images):
    features = []
    for image in tqdm(images):
        img = load_img(image,grayscale =  True )
        img = np.array(img)
        features.append(img)
    features = np.array(features)
    features = features.reshape(len(features),48,48,1)
test_features = extract_features(test['image'])
```

1. Import `tqdm` (Optional):

- `from tqdm.notebook import tqdm` (assuming this is the intended import): This line imports the `tqdm` library, which provides a progress bar for iterating through sequences. It's helpful for visualizing progress, especially when dealing with large datasets.

```
x_train = train_features/255.0
x_test = test_features/255.0
```

2. `extract_features` Function:

- This function takes a list of image paths (`images`) as input.
- It creates an empty list named `features` to store the extracted features.
- It iterates through each image path (`image`) in the provided list. The `tqdm` wrapper (optional) displays a progress bar during iteration.
- Inside the loop:
    - `img = load_img(image, grayscale=True)`: Loads the image using `load_img` (likely from Keras), converting it to grayscale if `grayscale=True`.
    - `img = np.array(img)`: Converts the loaded image to a NumPy array, which is a suitable format for machine learning models.
    - `features.append(img)`: Appends the processed image to the `features` list.
- After the loop:
    - `features = np.array(features)`: Converts the `features` list (now containing NumPy arrays) into a single NumPy array. This allows for easier manipulation and processing.
    - `features = features.reshape(len(features), 48, 48, 1)`: Reshapes the features array to a specific format that might be required by your machine learning model. Here:
        - `len(features)`: Sets the first dimension to the number of images (samples).
        - `48, 48`: Sets the second and third dimensions to 48x48, assuming your images are expected to be resized to this size.
        - `1`: Sets the fourth dimension to 1, indicating a single grayscale channel (since `grayscale=True` was used earlier).
- The function returns the processed features array.

3. Feature Extraction for Training and Testing Data:

- These lines call the `extract_features` function twice:

- The first call uses `train['image']` (list of image paths from the training DataFrame) to extract features and stores them in `train_features`.
- The second call uses `test['image']` (list of image paths from the testing DataFrame) to extract features and stores them in `test_features`.

4. Normalization (Optional):

These lines (assuming they're not commented out) perform normalization on the extracted features. They divide both `train_features` and `test_features` by 255.0 (assuming pixel values range from 0 to 255). This normalization helps some machine learning models converge better during training.

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
le.fit(train['label'])
```

```
▾ LabelEncoder
LabelEncoder()
```

```
y_train = le.transform(train['label'])
y_test = le.transform(test['label'])

y_train = to_categorical(y_train,num_classes = 7)
y_test = to_categorical(y_test,num_classes = 7)
```

1. Importing LabelEncoder

- This line imports the `LabelEncoder` class from the `sklearn.preprocessing` module. `LabelEncoder` is used to convert categorical labels (like text labels for image classes) into numerical representations suitable for machine learning models.

2. Creating a LabelEncoder Instance:

- This line creates an instance of the `LabelEncoder` class and assigns it to the variable `le`.

3. Fitting the LabelEncoder (Optional):

- This line (assuming it's not commented out) fits the `LabelEncoder` to the labels in the `'label'` column of the `train` DataFrame. Fitting involves learning the unique categories (class labels) present in the training data. This step might be optional depending on how `to_categorical` (explained later) is used.

4. Encoding Training Labels:

- This line uses the fitted `LabelEncoder` (`le`) to transform the labels in the `'label'` column of the `train` DataFrame. The `transform` method converts each label (category name) into a numerical representation based on the mapping learned during fitting. The result is stored in the variable `y_train`.

5. Encoding Testing Labels:

- Similar to `y_train`, this line transforms the labels in the `'label'` column of the `test` DataFrame using the same `LabelEncoder` (`le`). This ensures consistent encoding between training and testing data.

6. One-Hot Encoding (Optional):

• These lines perform one-hot encoding on the label-encoded data (`y_train` and `y_test`). One-hot encoding represents each category (class) as a fixed-length vector with all zeros except for the index corresponding to the class. Here:

- `to_categorical` is likely imported from `keras.utils` (assuming this code uses Keras).
- `num_classes=7`: This argument specifies that there are 7 possible classes (categories) in your image dataset.

• The result of `to_categorical` is a NumPy array where each row represents a sample (image) and each column represents a class. A 1 in a column indicates the corresponding class for that sample, and all other columns are 0

```
model = Sequential()
# convolutional layers
model.add(Conv2D(128, kernel_size=(3,3), activation='relu', input_shape=(48,48,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
```

```
model.add(Conv2D(256, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))

model.add(Conv2D(512, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))

model.add(Conv2D(512, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))

model.add(Flatten())
# fully connected layers
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.3))
# output layer
model.add(Dense(7, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x=x_train, y=y_train, batch_size=128, epochs=100, validation_data=(x_test, y_test))
```

```
model_json = model.to_json()
with open("emotiondetector.json",'w') as json_file:
    json_file.write(model_json)
model.save("emotiondetector.h5")
```

```
import cv2
import numpy as np
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Dropout, Flatten, MaxPooling2D

model = Sequential()
model.add(Conv2D(128, kernel_size=(3,3), activation='relu', input_shape=(48,48,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
model.add(Conv2D(256, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
model.add(Conv2D(512, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
model.add(Conv2D(512, kernel_size=(3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.4))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(7, activation='softmax'))
model.load_weights("facialemotionmodel.h5")


haar_file = cv2.data.haarcascades + 'haarcascade_frontalface_default.xml'


face_cascade = cv2.CascadeClassifier(haar_file)


def extract_features(image):
    feature = np.array(image)
    feature = feature.reshape(1, 48, 48, 1)
```

```python
    return feature / 255.0


labels = {0: 'angry', 1: 'disgust', 2: 'fear', 3: 'happy', 4: 'neutral', 5: 'sad', 6: 'surprise'}

# Open the webcam
webcam = cv2.VideoCapture(0)

while True:
    # Read a frame from the webcam
    ret, im = webcam.read()

    # Convert the frame to grayscale
    gray = cv2.cvtColor(im, cv2.COLOR_BGR2GRAY)

    # Detect faces in the grayscale frame
    faces = face_cascade.detectMultiScale(gray, 1.3, 5)

    try:
        # Iterate over the detected faces
        for (x, y, w, h) in faces:
            # Extract the region of interest (ROI) from the grayscale frame
            roi_gray = gray[y:y + h, x:x + w]

            # Draw a rectangle around the face in the original frame
            cv2.rectangle(im, (x, y), (x + w, y + h), (255, 0, 0), 2)


            roi_gray = cv2.resize(roi_gray, (48, 48))


            img = extract_features(roi_gray)


            pred = model.predict(img)


            prediction_label = labels[pred.argmax()]

            # Draw the predicted label on the original frame
            cv2.putText(im, '%s' % (prediction_label), (x - 10, y - 10), cv2.FONT_HERSHEY_COMPLEX_SMALL, 2,
(0, 0, 255))

        # Display the output frame
        cv2.imshow("Output", im)


        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
    except cv2.error:
        pass


webcam.release()
cv2.destroyAllWindows()
```

RESULT


original image is of happy
1/1 [==============================] - 0s 42ms/step
model prediction is happy
<matplotlib.image.AxesImage at 0x16ac00a8970>
0
10
20
30
40


original image is of disgust
1/1 [==============================] - 0s 57ms/step
model prediction is disgust
<matplotlib.image.AxesImage at 0x16abfef4d90>
0
10
20
30
40

CHAPTER - 5

OBSERVATION

Data Collection and Preprocessing:

- The chosen datasets for facial expression and text data cover a range of emotions and sentiments, providing a good foundation for training models.
- Preprocessing techniques such as grayscale conversion, resizing, normalization, and text cleaning are essential for preparing the data for model training.
- Techniques like data augmentation for facial images and text tokenization contribute to improving model robustness and performance.

Model Selection and Training:

- Convolutional Neural Networks (CNNs) are favored for facial expression recognition due to their ability to learn spatial features from images.
- Recurrent Neural Networks (RNNs) and transformers like BERT are suitable for analyzing sequential data like text, offering different advantages in capturing contextual information.
- Training processes involving data splitting, model optimization, and hyperparameter tuning are crucial for achieving optimal model performance.

Evaluation and Analysis:

- Metrics such as accuracy, precision, recall, and F1-score are appropriately chosen to assess model performance for both facial expression recognition and sentiment analysis.
- Comparative analysis between different models helps in identifying the most suitable architecture for the given data and project goals.
- Error analysis provides insights into the types of data points that models struggle with, guiding further refinement strategies.

Integration and Refinement:

- Integration of well-performing individual models for facial expression recognition and sentiment analysis into a single system facilitates a comprehensive understanding of emotions conveyed in both facial expressions and text.
- Ensemble learning techniques offer potential for further improving system performance through the combination of predictions from multiple models.
- Refinement strategies including experimentation with model architectures, data augmentation, and transfer learning contribute to enhancing the system's accuracy and generalizability.

Testing and Deployment:

- Testing the integrated system on unseen data ensures its robustness and applicability to real-world scenarios, validating its performance beyond the training data.
- Deployment options such as web applications or APIs provide avenues for wider usage and interaction with the developed system.

Reporting and Documentation:

- Comprehensive reporting and documentation are essential for summarizing the project, detailing methodologies, results, and future directions, thereby facilitating reproducibility and collaboration.

CHAPTER - 6

RESULT AND CONCLUSION

1. **Sentiment Analysis and Emotion Recognition System**

The project successfully developed a system for analyzing sentiment and recognizing emotions using both facial expressions and textual data. This accomplishment aligns with the project's initial goals of building a comprehensive system for understanding user sentiment and emotions through these combined modalities.

2. **Effective Deep Learning Models**

The project's utilization of deep learning models, such as Convolutional Neural Networks (CNNs) for facial expressions and Recurrent Neural Networks (RNNs) or transformers for sentiment analysis, demonstrated promising results. Evaluation metrics confirmed the effectiveness of these models in achieving accurate sentiment analysis and emotion recognition.

3. **Integrated Multimodal Analysis**

A key outcome of the project involved the successful integration of separate models for facial expression recognition and sentiment analysis. This integration allows the system to combine sentiment analysis results with the identified emotions from facial expressions, providing a more comprehensive picture of user sentiment and emotional state.

4. **Foundation for Future Advancements**

This project serves as a strong foundation for further exploration and refinement in the field of sentiment analysis and emotion recognition. The findings highlight the potential of deep learning models and multimodal analysis for achieving accurate results.

5. **Future Research Directions**

The project identified several promising avenues for future research. These include exploring multimodal fusion with additional modalities like vocal tone and body language, enabling real-time processing for practical applications, and considering cultural variations in emotional expression.

6. **Explainability and Security**

Future work should delve into Explainable AI (XAI) techniques to understand model predictions and mitigate potential biases. Additionally, addressing privacy and security concerns when handling sensitive data like facial expressions and text is crucial for responsible system development.

7. **Broader Applications**

By building upon these results and pursuing future research directions, this project opens doors for developing increasingly sophisticated systems with applications in various domains like healthcare, education, and human-computer interaction. The potential integration with Brain-Computer Interfaces (BCIs) presents exciting possibilities for future advancements.

CHAPTER - 7

RECOMMENDATION FOR FUTURE WORK

Building upon the current project, here are some exciting avenues for future exploration:

1. Multimodal Fusion:

- Explore techniques for combining information from facial expressions, text data, and potentially other modalities (e.g., vocal tone, body language) for a more comprehensive understanding of sentiment and emotions. This could involve developing deep learning models that can handle multimodal inputs or designing architectures that effectively combine the outputs from separate facial expression recognition and sentiment analysis models.

2. Real-Time Processing:

- Investigate methods for enabling real-time processing of facial expressions and text data. This could involve optimizing models for faster inference or utilizing specialized hardware for efficient computation. Real-time processing capabilities would open doors for applications like live customer service interactions, educational feedback systems, or even personalized advertising that adapts to a user's emotional state.

3. Cultural Considerations:

- Expand the project's scope to consider cultural variations in facial expressions and sentiment expression. Facial expressions and their interpretations can vary across cultures, and sentiment conveyed in text may depend on cultural norms and linguistic nuances. Develop models that can account for these variations or create culture-specific models for improved accuracy.

4. Explainable AI (XAI):

- Integrate Explainable AI (XAI) techniques to understand why the models make certain predictions, particularly when errors occur. This transparency can build trust in the system's outputs and help identify potential biases that might be present in the data or model architecture.

5. Privacy and Security:

- Address privacy and security concerns when dealing with facial expressions and text data, which can be considered sensitive information. Develop anonymization techniques or secure data collection and processing pipelines to ensure user privacy and data security.

6. Applications in Specific Domains:

- Explore applications of the system in specific domains like healthcare (detecting patient emotions during therapy sessions), education (providing personalized feedback based on student engagement), or human-computer interaction (creating emotionally intelligent chatbots). Adapt the system to the specific needs and data characteristics of these domains.

7. Integration with Brain-Computer Interfaces (BCIs):

- As Brain-Computer Interface (BCI) technology advances, explore the potential for combining sentiment analysis and emotion recognition with BCI data. This could enable the development of systems that can directly capture a user's emotional state and combine it with text analysis for a deeper understanding of human emotions.

These are just a few potential directions for future research. As technology progresses and our understanding of sentiment analysis and emotion recognition deepens, exciting possibilities will emerge for developing even more sophisticated and impactful systems.

CHAPTER - 8

REFERENCES

[1] Azevedo, A.: Data mining and knowledge discovery in databases. In: Advanced Methodologies and Technologies in Network Architecture, Mobile Computing, and Data Analytics, pp. 502–514. IGI Global, ??? (2019)

[2] Yalcin, N., DiPaola, S.: A computational model of empathy for interactive agents. Biologically inspired cognitive architectures 26, 20–25 (2018)

[3] Essa, I.A., Pentland, A.P.: Facial expression recognition using a dynamic model and motion energy. In: Proceedings of IEEE International Conference on Computer Vision, pp. 360–367 (1995). IEEE

[4] Pantic, M., Rothkrantz, L.J.M.: Automatic analysis of facial expressions: The state of the art. IEEE Transactions on pattern analysis and machine intelligence 22(12), 1424–1445 (2000)

[5] Cowie, P.A., Roberts, G.P.: Constraining slip rates and spacings for active normal faults. Journal of Structural Geology 23(12), 1901–1915 (2001)

[6] Valenti, R., Sebe, N., Gevers, T., et al.: Facial expression recognition: A fully integrated approach. In: Workshop on Visual and Multimedia Digital Libraries, vol. 10 (2007)

[7] Islam, J., Zhang, Y.: Visual sentiment analysis for social images using transfer learning approach. In: 2016 IEEE International Conferences on Big Data and Cloud Computing (BDCloud), Social Computing and Networking (Social-Com), Sustainable Computing and Communications (SustainCom)(BDCloud-

SocialCom-SustainCom), pp. 124–130 (2016). IEEE

[8] Severyn, A., Moschitti, A.: Twitter sentiment analysis with deep convolutional neural networks. In: Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval, pp. 959–962 (2015)

[9] You, Q., Luo, J., Jin, H., Yang, J.: Joint visual-textual sentiment analysis with deep neural networks. In: Proceedings of the 23rd ACM International Conference on Multimedia, pp. 1071–1074 (2015)

[10] Li, C., Xu, B., Wu, G., He, S., Tian, G., Hao, H.: Recursive deep learning for sentiment analysis over social data. In: 2014 IEEE/WIC/ACM International Joint Conferences on Web Intelligence (WI) and Intelligent Agent Technologies (IAT), vol. 2, pp. 180–185 (2014). IEEE

[11] Baecchi, C., Uricchio, T., Bertini, M., Del Bimbo, A.: A multimodal feature learning approach for sentiment analysis of social network multimedia. Multimedia Tools and Applications 75, 2507–2525 (2016)

[12] Yang, J., Sun, M., Sun, X.: Learning visual sentiment distributions via augmented conditional probability neural network. In: Proceedings of the AAAI Conference on Artificial Intelligence, vol. 31 (2017)

[13] Vadicamo, L., Carrara, F., Cimino, A., Cresci, S., Dell'Orletta, F., Falchi, F., Tesconi, M.: Cross-media learning for image sentiment analysis in the wild. In: Proceedings of the IEEE International Conference on Computer Vision Workshops, pp. 308–317 (2017)

[14] Siersdorfer, S., Minack, E., Deng, F., Hare, J.: Analyzing and predicting sentiment of images on the social web. In: Proceedings of the 18th ACM International Conference on Multimedia, pp. 715–718 (2010)

[15] Peng, K.-C., Chen, T., Sadovnik, A., Gallagher, A.C.: A mixed bag of emotions: Model, predict, and transfer emotion distributions. In: Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, pp. 860–868 (2015)

[16] Mikels, J.A., Fredrickson, B.L., Larkin, G.R., Lindberg, C.M., Maglio, S.J., Reuter-Lorenz, P.A.: Emotional category data on images from the international affective picture system. Behavior research methods 37, 626–630 (2005)

[17] Patel, K., Mehta, D., Mistry, C., Gupta, R., Tanwar, S., Kumar, N., Alazab, M.: Facial sentiment analysis using ai techniques: state-of-the-art, taxonomies, and challenges. IEEE access 8, 90495–90519 (2020)

[18] Torres, A.D., Yan, H., Aboutalebi, A.H., Das, A., Duan, L., Rad, P.: Patient facial emotion recognition and sentiment analysis using secure cloud with hardware acceleration. In: Computational Intelligence for Multimedia Big Data on the Cloud with Engineering Applications, pp. 61–89. Elsevier, ??? (2018)

[19] Kalla, D., Smith, N., Samaah, F., Polimetla, K.: Facial emotion and sentiment detection using convolutional neural network. Indian Journal of Artificial Intelligence Research (INDJAIR) 1(1), 1–13 (2021)

[20] Ali, M.F., Khatun, M., Turzo, N.A.: Facial emotion detection using neural network. the international journal of scientific and engineering research (2020)

[21] Wikarsa, L., Thahir, S.N.: A text mining application of emotion classifications of twitter's users using naive bayes method. In: 2015 1st International Conference on Wireless and Telematics (ICWT), pp. 1–6 (2015). IEEE

[22] Perikos, I., Hatzilygeroudis, I.: Recognizing emotions in text using ensemble of classifiers. Engineering Applications of Artificial Intelligence 51, 191–201 (2016)

[23] LeCompte, T., Chen, J.: Sentiment analysis of tweets including emoji data. In: 2017 International Conference on Computational Science and Computational Intelligence (CSCI), pp. 793–798 (2017). IEEE

[24] Kuˇsen, E., Cascavilla, G., Figl, K., Conti, M., Strembeck, M.: Identifying emotions in social media: comparison of word-emotion lexicons. In: 2017 5th International Conference on Future Internet of Things and Cloud Workshops (FiCloudW), pp. 132–137 (2017). IEEE

[25] Jain, V.K., Kumar, S., Fernandes, S.L.: Extraction of emotions from multilingual text using intelligent text processing and computational linguistics. Journal of computational science 21, 316–326 (2017)

[26] Badugu, S., Suhasini, M.: Emotion detection on twitter data using knowledge base approach. International Journal of Computer Applications 162(10) (2017)

[27] Kratzwald, B., Ili´c, S., Kraus, M., Feuerriegel, S., Prendinger, H.: Deep learning for affective computing: Text-based emotion recognition in decision support. Decision support systems 115, 24–35 (2018)

[28] Sailunaz, K., Alhajj, R.: Emotion and sentiment analysis from twitter text. Journal of computational science 36, 101003 (2019)

[29] Shrivastava, K., Kumar, S., Jain, D.K.: An effective approach for emotion detection in multimedia text data using sequence based convolutional neural network. Multimedia tools and applications 78, 29607–29639 (2019)

[30] Xu, G., Li, W., Liu, J.: A social emotion classification approach using multi-model fusion. Future Generation Computer Systems 102, 347–356 (2020)

[31] Adoma, A.F., Henry, N.-M., Chen, W.: Comparative analyses of bert, roberta, distilbert, and xlnet for text-based emotion recognition. In: 2020 17th International Computer Conference on Wavelet Active Media Technology and Information Processing (ICCWAMTIP), pp. 117–121 (2020). IEEE

[32] Chowanda, A., Sutoyo, R., Tanachutiwat, S., et al.: Exploring text-based emotions recognition machine learning techniques on social media conversation. Procedia Computer Science 179, 821–828 (2021)

[33] Dheeraj, K., Ramakrishnudu, T.: Negative emotions detection on online mental-health related patients texts using the deep learning with mha-bcnn model. Expert Systems with Applications 182, 115265 (2021)

[34] Santamaria-Granados, L., Munoz-Organero, M., Ramirez-Gonzalez, G., Abdulhay, E., Arunkumar, N.: Using deep convolutional neural network for emotion detection on a physiological signals dataset (amigos). IEEE Access 7, 57–67 (2018)